

## Project Overview

The LLM Evaluation Suite is a domain-agnostic framework for evaluating professional reasoning rather than isolated technical ability.

Traditional AI benchmarks typically measure whether a model can produce the correct answer to a predefined problem. While valuable, these benchmarks often reveal little about how an AI behaves in realistic professional environments where requirements are incomplete, ambiguity is expected, and sound judgment is more valuable than memorized knowledge.

This project evaluates the *process* behind a model's decisions. It measures how an AI gathers information, identifies missing requirements, determines when clarification is necessary, follows accepted professional standards, manages uncertainty, and justifies its reasoning before producing a solution.

The framework is intentionally designed to separate objective engineering principles ("baking") from subjective implementation preferences ("cooking"), allowing multiple correct solutions while rewarding sound professional practice instead of stylistic preference.

Although the initial evaluation scenarios focus on software engineering, the underlying methodology is designed to extend naturally to other disciplines—including healthcare, investigations, business operations, insurance, legal workflows, recovery services, and any profession where asking the right questions is often more important than giving immediate answers.

The ultimate objective is not to determine which AI is the most knowledgeable, but which AI consistently demonstrates the judgment, restraint, and reasoning expected of a trusted professional collaborator.